In [1]:
```python
import pandas as pd
```

In [2]:
```python
import numpy as np
```

In [3]:
```python
df= pd.read_csv('https://raw.githubusercontent.com/Lorddhaval/Dataset/main/Big%20Sales%20Data.csv')
```

In [4]:
```python
df.head()
```

Out[4]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales
0	FDT36	12.3	Low Fat	0.111448	Baking Goods	33.4874	OUT049	1999	Medium	Tier 1	Supermarket Type1	436.608721
1	FDT36	12.3	Low Fat	0.111904	Baking Goods	33.9874	OUT017	2007	Medium	Tier 2	Supermarket Type1	443.127721
2	FDT36	12.3	LF	0.111728	Baking Goods	33.9874	OUT018	2009	Medium	Tier 3	Supermarket Type2	564.598400
3	FDT36	12.3	Low Fat	0.000000	Baking Goods	34.3874	OUT019	1985	Small	Tier 1	Grocery Store	1719.370000
4	FDP12	9.8	Regular	0.045523	Baking Goods	35.0874	OUT017	2007	Medium	Tier 2	Supermarket Type1	352.874000

In [5]:
```python
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
 #   Column                     Non-Null Count  Dtype  
---  ------                     --------------  -----  
 0   Item_Identifier            14204 non-null  object 
 1   Item_Weight                11815 non-null  float64
 2   Item_Fat_Content           14204 non-null  object 
 3   Item_Visibility            14204 non-null  float64
 4   Item_Type                  14204 non-null  object 
 5   Item_MRP                   14204 non-null  float64
 6   Outlet_Identifier          14204 non-null  object 
 7   Outlet_Establishment_Year  14204 non-null  int64  
 8   Outlet_Size                14204 non-null  object 
 9   Outlet_Location_Type       14204 non-null  object 
 10  Outlet_Type                14204 non-null  object 
 11  Item_Outlet_Sales          14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

In [6]:
```python
df.describe()
```

Out[6]:

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	11815.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.788355	0.065953	141.004977	1997.830681	2185.836320
std	4.654126	0.051459	62.086938	8.371664	1827.479550
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	8.710000	0.027036	94.012000	1987.000000	922.135101
50%	12.500000	0.054021	142.247000	1999.000000	1768.287680
75%	16.750000	0.094037	185.855600	2004.000000	2988.110400
max	30.000000	0.328391	266.888400	2009.000000	31224.726950

In [7]:
```python
df.columns
```

Out[7]:
```
Index(['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibility',
       'Item_Type', 'Item_MRP', 'Outlet_Identifier',
       'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
       'Outlet_Type', 'Item_Outlet_Sales'],
      dtype='object')
```

In [8]:
```python
df['Item_Weight'].fillna(df.groupby(['Item_Type'])['Item_Weight'].transform('mean'),inplace =True)
```

In [9]:
```python
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
 #   Column                     Non-Null Count  Dtype  
---  ------                     --------------  -----  
 0   Item_Identifier            14204 non-null  object 
 1   Item_Weight                14204 non-null  float64
 2   Item_Fat_Content           14204 non-null  object 
 3   Item_Visibility            14204 non-null  float64
 4   Item_Type                  14204 non-null  object 
 5   Item_MRP                   14204 non-null  float64
 6   Outlet_Identifier          14204 non-null  object 
 7   Outlet_Establishment_Year  14204 non-null  int64  
 8   Outlet_Size                14204 non-null  object 
 9   Outlet_Location_Type       14204 non-null  object 
 10  Outlet_Type                14204 non-null  object 
 11  Item_Outlet_Sales          14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

In [10]:
```python
df.describe()
```

Out[10]:

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	14204.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.790642	0.065953	141.004977	1997.830681	2185.836320
std	4.251186	0.051459	62.086938	8.371664	1827.479550
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	9.300000	0.027036	94.012000	1987.000000	922.135101
50%	12.800000	0.054021	142.247000	1999.000000	1768.287680
75%	16.000000	0.094037	185.855600	2004.000000	2988.110400
max	30.000000	0.328391	266.888400	2009.000000	31224.726950

In [11]:
```python
import seaborn as sns
sns.pairplot(df)
```

Out[11]:
```
&lt;seaborn.axisgrid.PairGrid at 0x1a5df3cf5b0&gt;
```

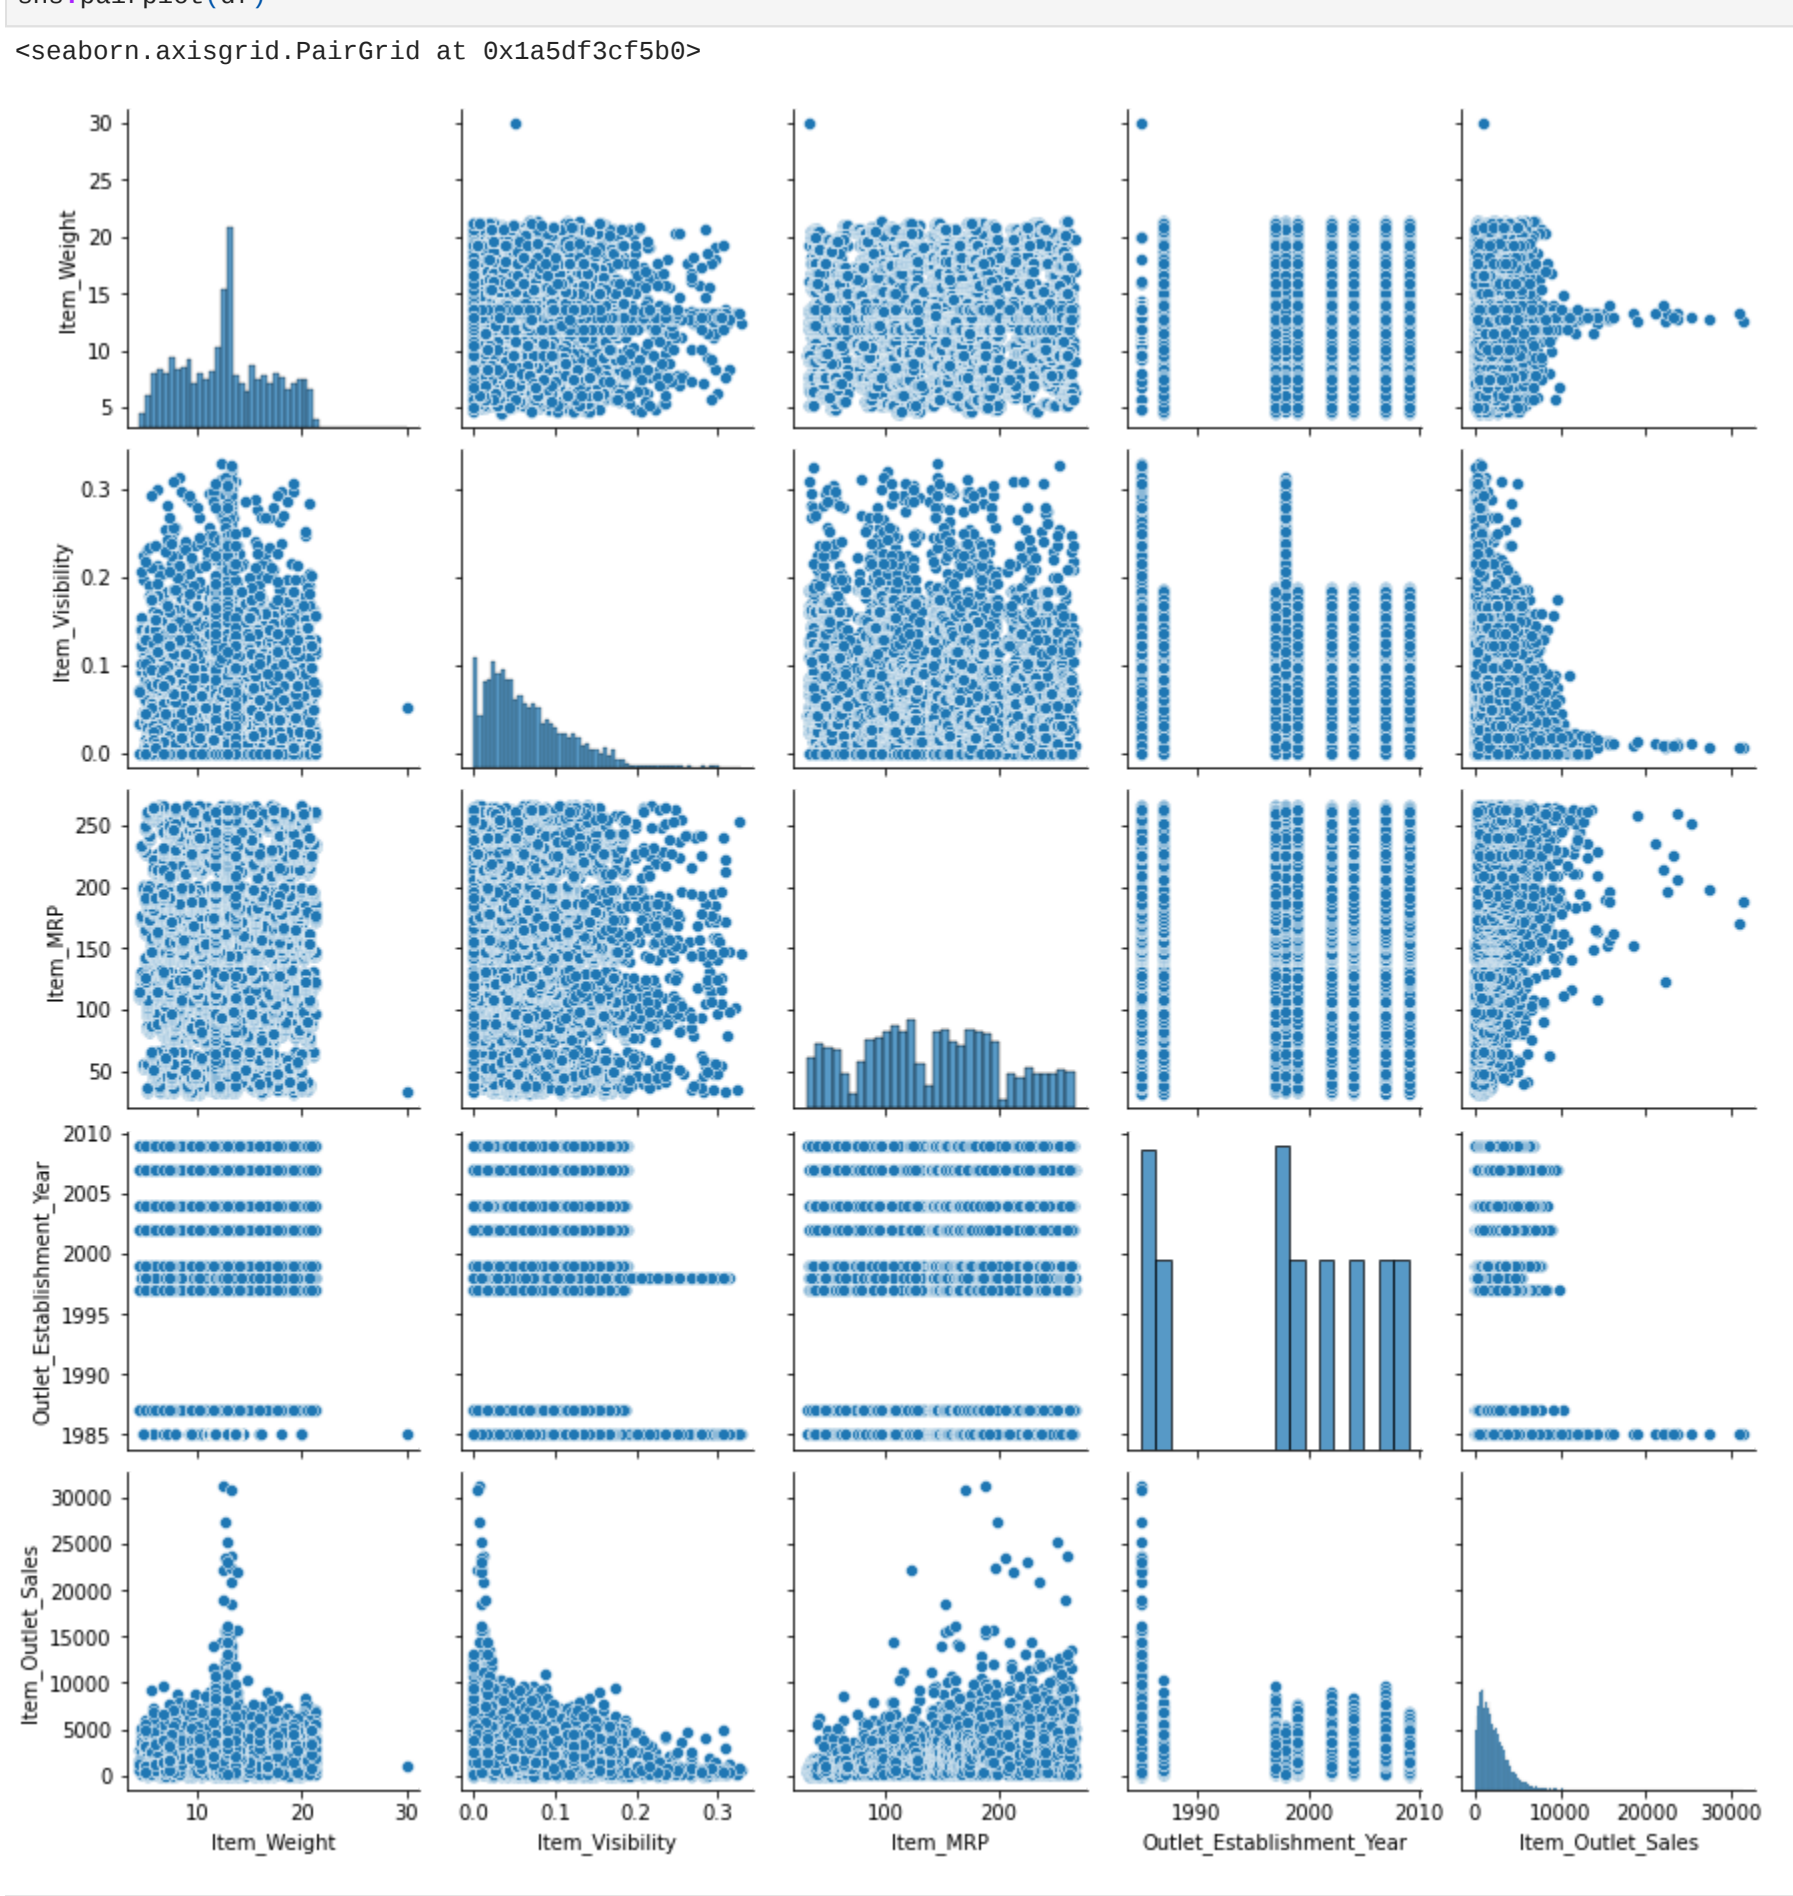

In [12]:
```python
df[['Item_Identifier']].value_counts()
```

Out[12]:
```
Item_Identifier
FDQ08              10
FDO24              10
FDQ19              10
FDQ28              10
FDQ31              10
                   ..
FDM52               7
FDM50               7
FDL50               7
FDM10               7
FDR51               7
Length: 1559, dtype: int64
```

In [13]:
```python
df[['Item_Fat_Content']].value_counts()
```

Out[13]:
```
Item_Fat_Content
Low Fat             8485
Regular             4824
LF                   522
reg                  195
low fat              178
dtype: int64
```

In [14]:
```python
df.replace({'Item_Fat_Content': {'LF':'Low Fat','reg':'Regular','low fat' :'Low Fat'}},inplace=True)
```

In [15]:
```python
df[['Item_Fat_Content']].value_counts()
```

Out[15]:
```
Item_Fat_Content
Low Fat             9185
Regular             5019
dtype: int64
```

In [16]:
```python
df.replace({'Item_Fat_Content': {'Low Fat':0,'Regular' :1}},inplace = True)
```

In [17]:
```python
df[['Item_Type']].value_counts()
```

Out[17]:
```
Item_Type
Fruits and Vegetables    2013
Snack Foods              1989
Household                1548
Frozen Foods             1426
Dairy                    1136
Baking Goods             1086
Canned                   1084
Health and Hygiene        858
Meat                      736
Soft Drinks               726
Breads                    416
Hard Drinks               362
Others                    280
Starchy Foods             269
Breakfast                 186
Seafood                    89
dtype: int64
```

In [18]:
```python
df.replace({'Item_Type':{'Fruits and Vegetables':0, 'Snack Foods':0, 'Household':1,
'Frozen Foods': 0, 'Dairy': 0, 'Baking Goods' : 0, 'Canned': 0, 'Health and Hygiene': 1, 'Meat': 0, 'Soft Drinks': 0,
'Breads': 0, 'Hard Drinks' : 0, 'Others' : 2, 'Starchy Foods': 0, 'Breakfast': 0, 'Seafood' : 0 }},inplace=True)
```

In [19]:
```python
df[['Item_Type']].value_counts()
```

Out[19]:
```
Item_Type
0            11518
1             2406
2              280
dtype: int64
```

In [20]:
```python
df[['Item_Type']].info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 1 columns):
 #   Column     Non-Null Count  Dtype
---  ------     --------------  -----
 0   Item_Type  14204 non-null  int64
dtypes: int64(1)
memory usage: 111.1 KB
```

In [21]:
```python
df[['Outlet_Identifier']].value_counts()
```

Out[21]:
```
Outlet_Identifier
OUT027               1559
OUT013               1553
OUT035               1550
OUT046               1550
OUT049               1550
OUT045               1548
OUT018               1546
OUT017               1543
OUT010                925
OUT019                880
dtype: int64
```

In [22]:
```python
df.replace({'Outlet_Identifier': { 'OUT027': 0, 'OUT013': 1,

'OUT049': 2, 'OUT046': 3, 'OUT035' : 4,

'OUT045':5, 'OUT018' : 6,

'OUT017': 7, 'OUT010': 8, 'OUT019' : 9,

}},inplace=True)
```

In [23]:
```python
df[['Outlet_Identifier']].value_counts()
```

Out[23]:
```
Outlet_Identifier
0                    1559
1                    1553
2                    1550
3                    1550
4                    1550
5                    1548
6                    1546
7                    1543
8                     925
9                     880
dtype: int64
```

In [24]:
```python
df[['Outlet_Size']].value_counts()
```

Out[24]:
```
Outlet_Size
Medium         7122
Small          5529
High           1553
dtype: int64
```

In [25]:
```python
df.replace({'Outlet_Size':{'Small':0, 'Medium':1,'High':2}},inplace=True)
```

In [26]:
```python
df[['Outlet_Location_Type']].value_counts()
```

Out[26]:
```
Outlet_Location_Type
Tier 3                  5583
Tier 2                  4641
Tier 1                  3980
dtype: int64
```

In [27]:
```python
df.replace({'Outlet_Location_Type':{'Tier 1':0, 'Tier 2':1,'Tier 3':2}},inplace=True)
```

In [28]:
```python
df[['Outlet_Location_Type']].value_counts()
```

Out[28]:
```
Outlet_Location_Type
2                       5583
1                       4641
0                       3980
dtype: int64
```

In [29]:
```python
df[['Outlet_Type']].value_counts()
```

Out[29]:
```
Outlet_Type
Supermarket Type1    9294
Grocery Store        1805
Supermarket Type3    1559
Supermarket Type2    1546
dtype: int64
```

In [30]:
```python
df.replace({'Outlet_Type':{'Grocery Store':0, 'Supermarket Type1':1,'Supermarket Type2':2,'Supermarket Type3':3}},inplace=True)
```

In [31]:
```python
df[['Outlet_Type']].value_counts()
```

Out[31]:
```
Outlet_Type
1              9294
0              1805
3              1559
2              1546
dtype: int64
```

In [32]:
```python
df.head()
```

Out[32]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales
0	FDT36	12.3	0	0.111448	0	33.4874	2	1999	1	0	1	436.608721
1	FDT36	12.3	0	0.111904	0	33.9874	7	2007	1	1	1	443.127721
2	FDT36	12.3	0	0.111728	0	33.9874	6	2009	1	2	2	564.598400
3	FDT36	12.3	0	0.000000	0	34.3874	9	1985	0	0	0	1719.370000
4	FDP12	9.8	1	0.045523	0	35.0874	7	2007	1	1	1	352.874000

In [33]:
```python
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
 #   Column                     Non-Null Count  Dtype  
---  ------                     --------------  -----  
 0   Item_Identifier            14204 non-null  object 
 1   Item_Weight                14204 non-null  float64
 2   Item_Fat_Content           14204 non-null  int64  
 3   Item_Visibility            14204 non-null  float64
 4   Item_Type                  14204 non-null  int64  
 5   Item_MRP                   14204 non-null  float64
 6   Outlet_Identifier          14204 non-null  int64  
 7   Outlet_Establishment_Year  14204 non-null  int64  
 8   Outlet_Size                14204 non-null  int64  
 9   Outlet_Location_Type       14204 non-null  int64  
 10  Outlet_Type                14204 non-null  int64  
 11  Item_Outlet_Sales          14204 non-null  float64
dtypes: float64(4), int64(7), object(1)
memory usage: 1.3+ MB
```

In [34]:
```python
df.shape
```

Out[34]:
```
(14204, 12)
```

In [35]:
```python
y = df['Item_Outlet_Sales']
```

In [36]:
```python
y.shape
```

Out[36]:
```
(14204,)
```

In [37]:
```python
y
```

Out[37]:
```
0         436.608721
1         443.127721
2         564.598400
3        1719.370000
4         352.874000
            ...
14199    4984.178800
14200    2885.577200
14201    2885.577200
14202    3803.676434
14203    3644.354765
Name: Item_Outlet_Sales, Length: 14204, dtype: float64
```

In [38]:
```python
X=df[['Item_Weight', 'Item_Fat_Content', 'Item_Visibility', 'Item_Type', 'Item_MRP',
      'Outlet_Identifier', 'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type','Outlet_Type']]
```

In [39]:
```python
X = df.drop(['Item_Identifier','Item_Outlet_Sales'],axis = 1)
```

In [40]:
```python
X.shape
```

Out[40]:
```
(14204, 10)
```

In [41]:
```python
X
```

Out[41]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type
0	12.300000	0	0.111448	0	33.4874	2	1999	1	0	1
1	12.300000	0	0.111904	0	33.9874	7	2007	1	1	1
2	12.300000	0	0.111728	0	33.9874	6	2009	1	2	2
3	12.300000	0	0.000000	0	34.3874	9	1985	0	0	0
4	9.800000	1	0.045523	0	35.0874	7	2007	1	1	1
...	...	...	...	...	...	...	...	...	...	...
14199	12.800000	0	0.069606	0	261.9252	4	2004	0	1	1
14200	12.800000	0	0.070013	0	262.8252	7	2007	1	1	1
14201	12.800000	0	0.069561	0	263.0252	1	1987	2	2	1
14202	13.659758	0	0.069282	0	263.5252	0	1985	1	2	3
14203	12.800000	0	0.069727	0	263.6252	2	1999	1	0	1

14204 rows × 10 columns

In [42]:
```python
from sklearn.preprocessing import StandardScaler
```

In [43]:
```python
sc = StandardScaler()
```

In [44]:
```python
X_std = df[['Item_Weight','Item_Visibility','Item_MRP','Outlet_Establishment_Year']]
```

In [45]:
```python
X_std = sc.fit_transform(X_std)
```

In [46]:
```python
X_std
```

Out[46]:
```
array([[-0.11541705,  0.88413635, -1.73178716,  0.13968068],
       [-0.11541705,  0.89300616, -1.72373366,  1.09531886],
       [-0.11541705,  0.88958331, -1.72373366,  1.3342284 ],
       ...,
       [ 0.00220132,  0.07011952,  1.96538148, -1.29377659],
       [ 0.20444792,  0.06469366,  1.97343499, -1.53268614],
       [ 0.00220132,  0.07334891,  1.97504569,  0.13968068]])
```

In [47]:
```python
X[['Item_Weight', 'Item_Visibility', 'Item_MRP', 'Outlet_Establishment_Year' ]] = pd.DataFrame(X_std, columns = [['Item_Weight', 'Item_Visibility','Item_MRP','Outlet_Establishment_Y
```

In [48]:
```python
X
```

Out[48]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location_Type	Outlet_Type
0	-0.115417	0	0.884136	0	-1.731787	2	0.139681	1	0	1
1	-0.115417	0	0.893006	0	-1.723734	7	1.095319	1	1	1
2	-0.115417	0	0.889583	0	-1.723734	6	1.334228	1	2	2
3	-0.115417	0	-1.281712	0	-1.717291	9	-1.532686	0	0	0
4	-0.703509	1	-0.397031	0	-1.706016	7	1.095319	1	1	1
...	...	...	...	...	...	...	...	...	...	...
14199	0.002201	0	0.070990	0	1.947664	4	0.736955	0	1	1
14200	0.002201	0	0.078898	0	1.962160	7	1.095319	1	1	1
14201	0.002201	0	0.070120	0	1.965381	1	-1.293777	2	2	1
14202	0.204448	0	0.064694	0	1.973435	0	-1.532686	1	2	3
14203	0.002201	0	0.073349	0	1.975046	2	0.139681	1	0	1

14204 rows × 10 columns

In [49]:
```python
from sklearn.model_selection import train_test_split
```

In [50]:
```python
X_train, X_test, y_train, y_test = train_test_split(X,y , test_size =0.1, random_state= 72529 )
```

In [51]:
```python
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[51]:
```
((12783, 10), (1421, 10), (12783,), (1421,))
```

In [52]:
```python
from sklearn.ensemble import RandomForestRegressor
```

In [53]:
```python
rfr = RandomForestRegressor(random_state=72529)
```

In [54]:
```python
rfr.fit(X_train, y_train)
```

Out[54]:
```
▾ RandomForestRegressor
RandomForestRegressor(random_state=72529)
```

In [55]:
```python
y_pred = rfr.predict(X_test)
```

In [56]:
```python
y_pred.shape
```

Out[56]:
```
(1421,)
```

In [57]:
```python
y_pred
```

Out[57]:
```
array([1513.89681418, 1796.47737123, 1896.58147176, ..., 3678.89208551,
        745.11958968, 3592.32874778])
```

In [58]:
```python
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

In [59]:
```python
mean_squared_error(y_test, y_pred)
```

Out[59]:
```
1480354.1533997788
```

In [60]:
```python
mean_absolute_error(y_test, y_pred)
```

Out[60]:
```
810.2512914480754
```

In [61]:
```python
r2_score(y_test, y_pred)
```

Out[61]:
```
0.5050364101543503
```

In [62]:
```python
import matplotlib.pyplot as plt
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Price Vs Predicted Price")
plt.show()
```

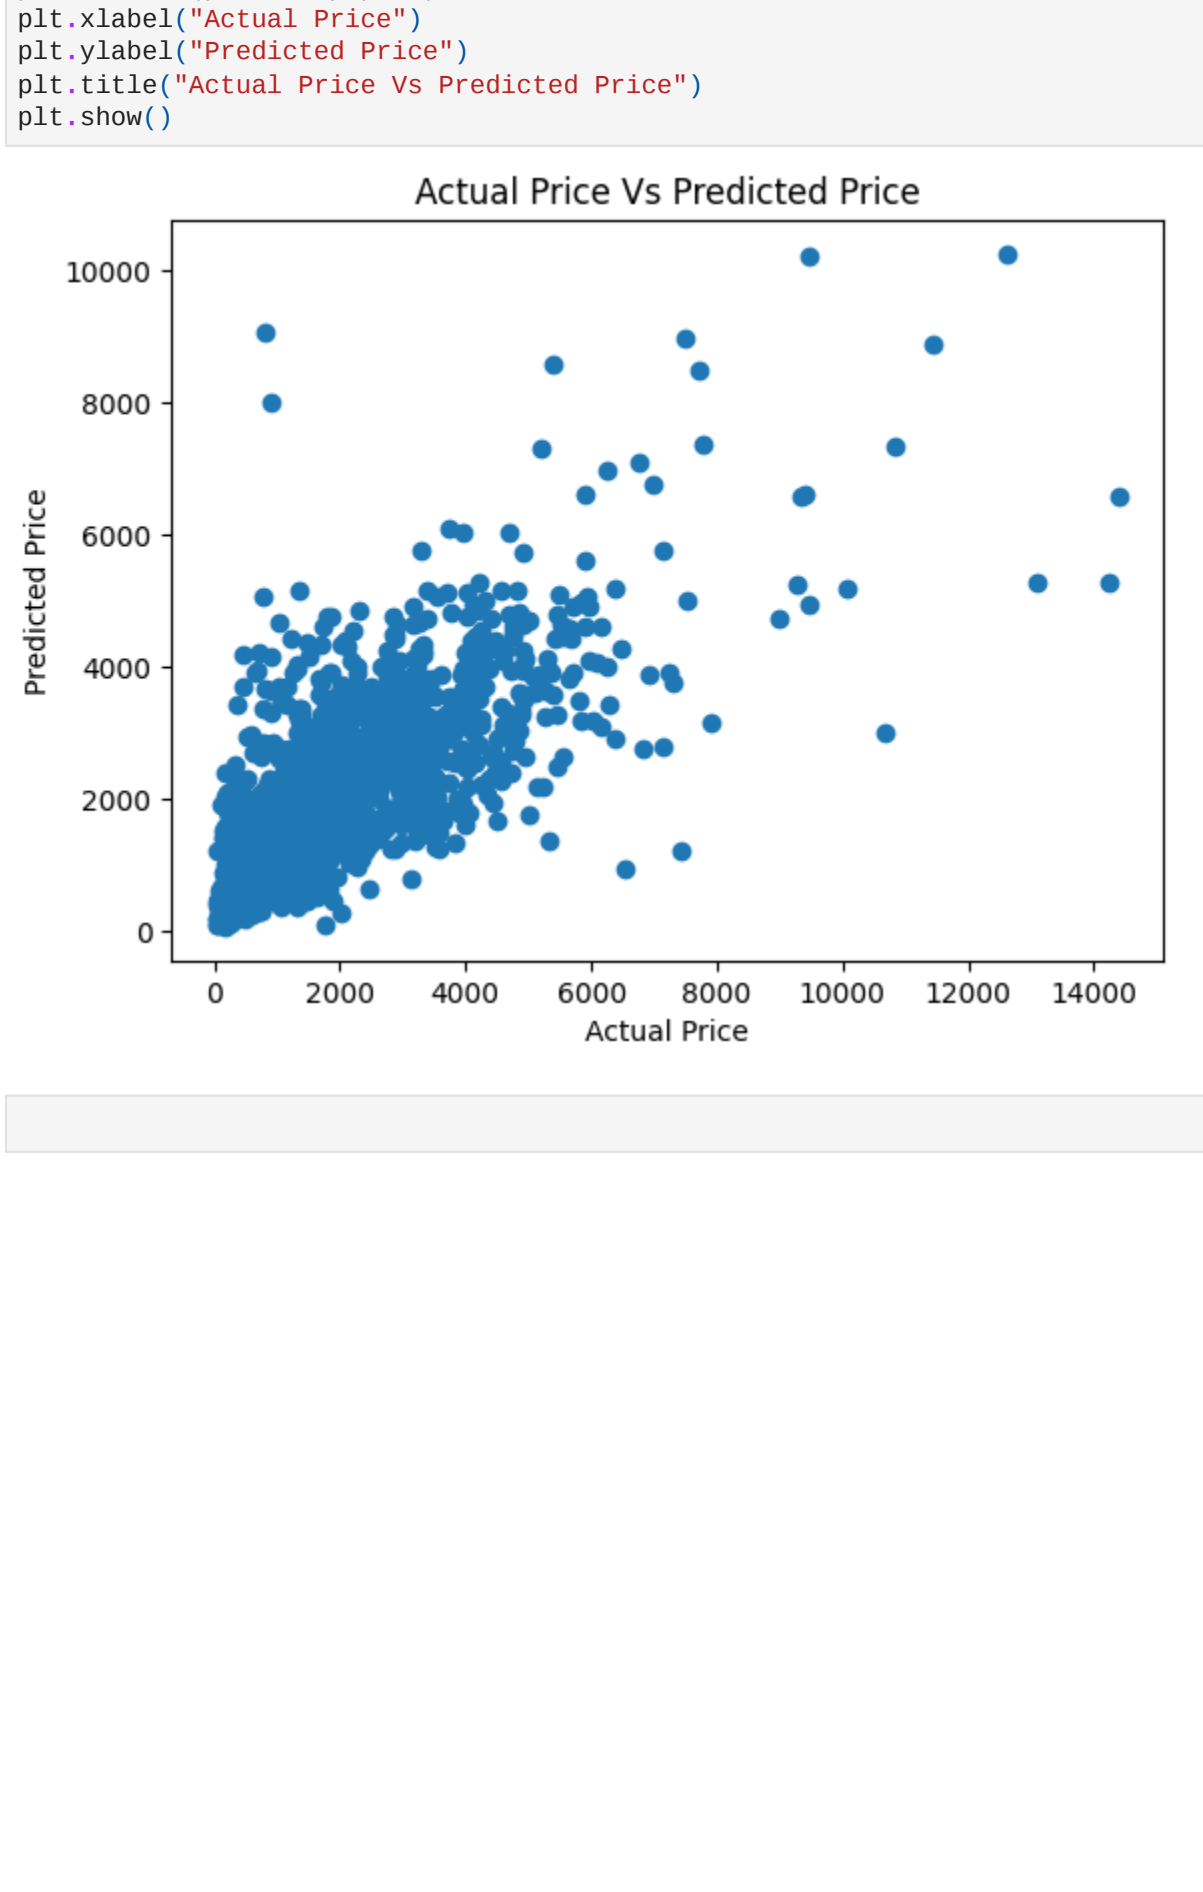

In [ ]: