

```
In [49]: import pandas as pd

In [2]: import numpy as np

In [3]: df = pd.read_csv(r'https://github.com/YBI-Foundation/Dataset/raw/main/Bike%20Prices.csv')

In [4]: df.head()

Out[4]:
```

| | Brand | Model | Selling_Price | Year | Seller_Type | Owner | KM_Driven | Ex_Showroom_Price |
|---|-------|--------------------|---------------|------|-------------|-----------|-----------|-------------------|
| 0 | TVS | TVS XL 100 | 30000 | 2017 | Individual | 1st owner | 8000 | 30490.0 |
| 1 | Bajaj | Bajaj ct 100 | 18000 | 2017 | Individual | 1st owner | 35000 | 32000.0 |
| 2 | Yo | Yo Style | 20000 | 2011 | Individual | 1st owner | 10000 | 37675.0 |
| 3 | Bajaj | Bajaj Discover 100 | 25000 | 2010 | Individual | 1st owner | 43000 | 42859.0 |
| 4 | Bajaj | Bajaj Discover 100 | 24999 | 2012 | Individual | 2nd owner | 35000 | 42859.0 |

```
In [5]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1061 entries, 0 to 1060
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                  1061 non-null   object
1   Model                  1061 non-null   object
2   Selling_Price          1061 non-null   int64
3   Year                   1061 non-null   int64
4   Seller_Type            1061 non-null   object
5   Owner                  1061 non-null   object
6   KM_Driven              1061 non-null   int64
7   Ex_Showroom_Price     626 non-null    float64
dtypes: float64(1), int64(3), object(4)
memory usage: 66.4+ KB

In [6]: df = df.dropna()

In [7]: df.describe()

Out[7]:
```

| | Selling_Price | Year | KM_Driven | Ex_Showroom_Price |
|-------|---------------|-------------|---------------|-------------------|
| count | 626.000000 | 626.000000 | 626.000000 | 6.260000e+02 |
| mean | 59445.164537 | 2014.800319 | 32671.576677 | 8.795871e+04 |
| std | 59904.350888 | 3.018885 | 45479.661039 | 7.749659e+04 |
| min | 6000.000000 | 2001.000000 | 380.000000 | 3.049000e+04 |
| 25% | 30000.000000 | 2013.000000 | 13031.250000 | 5.485200e+04 |
| 50% | 45000.000000 | 2015.000000 | 25000.000000 | 7.275250e+04 |
| 75% | 65000.000000 | 2017.000000 | 40000.000000 | 8.703150e+04 |
| max | 760000.000000 | 2020.000000 | 585659.000000 | 1.278000e+06 |

```
In [8]: df[['Brand']].value_counts()

Out[8]:
```

| Brand | |
|----------|-----|
| Honda | 179 |
| Bajaj | 143 |
| Hero | 108 |
| Yamaha | 94 |
| Royal | 48 |
| TVS | 23 |
| Suzuki | 18 |
| KTM | 6 |
| Mahindra | 6 |
| Kawasaki | 4 |
| UM | 3 |
| Activa | 3 |
| Harley | 2 |
| Vespa | 2 |
| BMW | 1 |
| Hyosung | 1 |
| Benelli | 1 |
| Yo | 1 |

```
In [9]: df[['Model']].value_counts()

Out[9]:
```

| Model | |
|--|----|
| Honda Activa [2000-2015] | 23 |
| Honda CB Hornet 160R | 22 |
| Bajaj Pulsar 180 | 20 |
| Yamaha FZ S V 2.0 | 16 |
| Bajaj Discover 125 | 16 |
| Royal Enfield Thunderbird 500 | 1 |
| Royal Enfield Continental GT [2013 - 2018] | 1 |
| Royal Enfield Classic Stealth Black | 1 |
| Royal Enfield Classic Squadron Blue | 1 |
| Yo Style | 1 |

```
In [10]: df[['Seller_Type']].value_counts()

Out[10]:
```

| Seller_Type | |
|-------------|-----|
| Individual | 623 |
| Dealer | 3 |

```
In [11]: df[['Owner']].value_counts()

Out[11]:
```

| Owner | |
|-----------|-----|
| 1st owner | 556 |
| 2nd owner | 66 |
| 3rd owner | 3 |
| 4th owner | 1 |

```
In [12]: df.columns

Out[12]:
```

| Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type', 'Owner', 'KM_Driven', 'Ex_Showroom_Price'], dtype='object') |
|--|
|--|

```
In [13]: df.shape

Out[13]:
```

| (626, 8) |
|----------|
|----------|

```
In [14]: df.replace({'Seller_Type':{'Individual':0,'Dealer':1}},inplace=True)

In [16]: df.replace({'Owner':{'1st owner':0,'2nd owner':1,'3rd owner':2,'4th owner':3}},inplace=True)

In [17]: #X = pd.get_dummies(X,columns =['Seller_Type0', 'Owner'],drop_first=True)

In [18]: y = df['Selling_Price']

In [19]: y.shape

Out[19]:
```

| (626,) |
|---------|
|---------|

```
In [20]: y

Out[20]:
```

| y | |
|-----|--------|
| 0 | 30000 |
| 1 | 18000 |
| 2 | 20000 |
| 3 | 25000 |
| 4 | 24999 |
| ... | ... |
| 621 | 330000 |
| 622 | 300000 |
| 623 | 425000 |
| 624 | 760000 |
| 625 | 750000 |

```
Name: Selling_Price, Length: 626, dtype: int64

In [21]: x = df[['Year', 'Seller_Type', 'Owner', 'KM_Driven', 'Ex_Showroom_Price']]

In [22]: #X df.drop (['Brand', 'Model', 'Selling_Price'],axis=1)

In [23]: x.shape

Out[23]:
```

| (626, 5) |
|----------|
|----------|

```
In [26]: from sklearn.model_selection import train_test_split

In [27]: x_train,x_test,y_train,y_test = train_test_split(x,y, test_size = 0.3, random_state=72529)

In [28]: x_train.shape,x_test.shape,y_train.shape,y_test.shape

Out[28]:
```

| ((438, 5), (188, 5), (438,), (188,)) |
|--------------------------------------|
|--------------------------------------|

```
In [29]: from sklearn.linear_model import LinearRegression

In [30]: lr = LinearRegression()

In [32]: lr.fit(x_train,y_train)

Out[32]:
```

| LinearRegression |
|--------------------|
| LinearRegression() |

```
In [33]: y_pred = lr.predict(x_test)

In [34]: y_pred.shape

Out[34]:
```

| (188,) |
|---------|
|---------|

```
In [35]: y_pred

Out[35]:
```

| array([138240.29654226, 35307.27629811, 41392.3146781 , 59696.08614052, 78589.64050907, 66330.33181893, 71450.75913796, 128741.92385458, 31301.4907403 , 16365.36217535, 52645.82622187, 34162.66438303, 33109.28322892, 29285.80469509, 22481.73317954, 18623.14544884, 46071.59537129, 43382.26599436, 49240.05425507, 52547.07580469, 63756.04007575, 49201.36726671, 60476.56587229, 46568.1904032 , 51727.28970532, 57675.94515841, 52904.87730434, 22154.32221548, 64738.47935315, 69961.82756377, 76237.03364442, 10109.81723728, 7978.29684737, 50290.33945634, 39002.28085522, 37636.24973153, 22214.51509801, 25919.18719911, 49681.05425329, 54478.5440337 , 134286.41450486, 39465.59348847, 18856.93520069, 29395.67326193, 52280.49539301, 54281.39627374, 23871.16128248, 22768.36144976, 51847.70627002, 74002.44137979, 60726.08485228, 55554.80090319, 84302.58855953, 78761.17925924, 13940.70161622, 116019.0995354, 14081.7149846, 44665.67959531, 131436.62182645, 22829.68163124, 8188.18124434, 98989.0543751 , 73329.16019507, 59492.70884178, 55000.52645284, 24443.01694916, 20675.98010482, 55621.8738598 , 24746.73279912, 65067.00272016, 24599.65499752, 50617.16305576, 47910.83100397, 49125.22125809, 54742.72180424, 42508.08197345, 38727.07477456, 48829.02962818, 64112.04371617, 53584.01982098, 54943.33533099, 67692.86373031, 126519.60362767, 26501.93016865, 67856.9440369 , 36218.20246254, 34696.88795136, 51491.4898429 , 86313.65522794, 42211.37363751, 45322.89571985, 47799.97285654, 102932.03791725, 32196.93603661, 45716.09120145, 69711.00966411, 16931.96292963, 20489.88095332, 28963.45154368, 44739.34616288, 63468.42147204, 66619.70418353, 24182.60691813, 31899.66302198, 33228.48506905, 55251.34662947, 21012.87110226, 38759.45196052, 60393.40507015, 54281.22548207, 52922.30019798, 64989.86871562, 118620.02866249, 52781.43797999, 50481.02485062, 54600.50644429, 110740.52754662, 48604.39480631, 140992.5623721 , 91457.62756442, 37668.69451781, 73385.8202705 , 251938.90006593, 115841.02501263, 136531.63426807, 28238.8208474 , 52842.56315046, 63314.77014161, 31188.37127546, 37092.05966479, 78643.45979421, 50279.40061888, 38977.99796575, 31588.58401208, 36994.47781951, 38419.49150795, 31414.013563 , 130271.76136431, 35704.39301389, 54191.15197736, 51895.22928438, 31043.80123759, 32558.17987778, 22186.69940144, 15425.81074276, 32115.77810923, 41215.81820956, 72011.78109481, 126282.35419302, 20390.2663526 , 55400.39036476, 50864.94573893, 87770.89530386, 23629.54522686, 131453.06572154, 57428.55542153, -2208.22086481, 58451.70345829, 112323.27244037, 51881.06551377, 40427.81846939, 94861.96975418, 61164.10657223, 68346.70230126, 52331.21702063, 27752.47525724, 73401.31807354, 36696.11481079, 36070.20279696, 25901.3679022 , 58724.9284498 , 55704.85706306, 60688.55113162, 59359.07183992, 46924.33944876, 50465.30096557, 44875.83392819, 41411.15345077, 59909.89024773, 51345.79150608, 34454.61844574, 71609.40734916, 154935.58973255, 127381.82858877, 60007.02109561, 69793.83455607, 29241.50034586, 54255.60205173]) |
|---|
|---|

```
In [36]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

In [37]: mean_squared_error(y_test,y_pred)

Out[37]:
```

| 152201375.067118 |
|------------------|
|------------------|

```
In [38]: mean_absolute_error(y_test,y_pred)

Out[38]:
```

| 8925.457899521194 |
|-------------------|
|-------------------|

```
In [39]: r2_score(y_test,y_pred)

Out[39]:
```

| 0.9003083084246579 |
|--------------------|
|--------------------|

```
In [40]: import matplotlib.pyplot as plt

In [42]: plt.scatter(y_test,y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Prediction Prices')
plt.title("Actual Prices vs Prediction Prices")
plt.show()
```